


## PA1 Report - Matrix Multiplication

### Q1 Results

#### Q1a

Matrix Size	Peak GF achieved in our Implementation
32	13.295
64	17.445
128	19.225
256	20.895
511	21.745
512	21.78
513	21.735
1023	21.63
1024	21.54
1025	21.615
2047	22.035
2048	21.78

#### Q1b



### Q2 Analysis

#### Q2a Working of Program

The following sequence of steps describe our code workflow:

1. **bl\_config.h** file defines the values for  $M_c$ ,  $N_c$ ,  $K_c$ ,  $M_r$ , and  $N_r$  parameters that are going to be used for our matrix multiplication. This file also declares `BL_MICRO_KERNEL` variable to `dgemm_14x4x4`, where `dgemm_14x4x4` is the name of the microkernel that we have implemented for matrix multiplication.
2. **bl\_dgemm\_kernel.h** file declares all the different size microkernels ( $16 \times 4 \times 4$ ,  $14 \times 4 \times 4$ ,  $12 \times 4 \times 4$ ,  $10 \times 4 \times 4$ ,  $8 \times 4 \times 4$ ,  $6 \times 4 \times 4$ ,  $4 \times 4 \times 4$ ,  $2 \times 4 \times 4$ ) that we have implemented as a part of this assignment.
3. **bl\_dgemm\_ukr.c** file contains definitions of all the different size microkernels that we have implemented as a part of this assignment. Function name `dgemm_14x4x4` corresponds to the microkernel that performs outer product of  $14 \times 4$  ( $M_r \times K_c$ ) subpanel of matrix A with a  $4 \times 4$  ( $K_c \times N_r$ ) subpanel of matrix B.
4. **my\_dgemm.c** file contains the code for actual matrix multiplication. `square_dgemm()` function invokes `bl_dgemm()` function which contains three loops followed by invocation of `bl_macro_kernel()`.
  - a. The 5th loop around microkernel (as commented in code) partitions matrix A and C into submatrices of size  $M_c \times n$ .
  - b. The 4th loop around microkernel (as commented in code) further partitions A into panels of size  $M_c \times K_c$ , and partitions matrix B into submatrix of size  $K_c \times n$ .
    - i. This loop also contains a loop to pack subpanels of size  $M_r \times K_c$  into each panel of size  $M_c \times K_c$ .
    - ii. Each  $M_r \times K_c$  subpanel is accessed in column major order and elements are saved in that order in an array named `packA`.
  - c. The 3rd loop around microkernel (as commented in code) further partitions B into panels of size  $K_c \times N_c$ .
    - i. This loop also contains a loop to pack subpanels of size  $K_c \times N_r$  into each panel of size  $K_c \times N_c$ .
    - ii. Each  $K_c \times N_r$  subpanel is accessed in row major order and elements are saved in that order in an array named `packB`.
  - d. The `bl_macro_kernel()` contains the 2nd and 1st loop around microkernel (as commented in code) that passes a  $M_r \times K_c$  subpanel of A and  $K_c \times N_r$  subpanel of B to microkernel.
  - e. Finally, the `dgemm_14x4x4()` microkernel reads 14 values stored in  $M_r \times N_r$  subpanel in matrix C, and one by one performs an outer product of a column of  $14 \times 4$  ( $M_r \times K_c$ ) subpanel of matrix A with a row of  $4 \times 4$  ( $K_c \times N_r$ ) subpanel of matrix B. It makes use of SVE instructions to multiply an element of matrix A with four elements in row of matrix B in a single instruction.

## Q2b Development Process


This section describes the seven optimizations that we tried for this assignment. The first four optimizations yielded us performance gain while the rest three optimizations did not give us any performance gain.

### Optimization 1: Implementing a microkernel using scalable vector registers and SVE instructions

Commit Message	Commit Link
Finished packing for naive microkernel	<a href="https://github.com/cse260-fa22/pa1-hgondali-shanant/commit/9b6ba86c74ee9f0e2aebd94163ff12a6cdaf6edf">https://github.com/cse260-fa22/pa1-hgondali-shanant/commit/9b6ba86c74ee9f0e2aebd94163ff12a6cdaf6edf</a>
Added 4x4 microkernel	<a href="https://github.com/cse260-fa22/pa1-hgondali-shanant/commit/811619117ae7f8e82d2380070d7ecd5e5607e624">https://github.com/cse260-fa22/pa1-hgondali-shanant/commit/811619117ae7f8e82d2380070d7ecd5e5607e624</a>

**Methodology and Reasoning:** We implemented a  $4 \times 4$  ( $M_r \times N_r$ ) microkernel that performs matrix multiplication of  $M_r \times K_c$  subpanel of A with  $K_c \times N_r$  subpanel of B using SIMD-based SVE instructions. This

microkernel is compared against a naive microkernel that performs the same computation using SISD-based instructions. These microkernels were used to perform matrix multiplications for different matrix sizes and their results are compared.



**Results Analysis:** In the case of SISD instructions based microkernel, a single instruction can perform a single multiplication operation for a single set of operands. On the other hand, SVE instructions adopt a SIMD approach where four multiplication operation can be performed at the same time for four independent sets of operands. Thus, SVE instruction based microkernel can perform the same number of operations in significantly lower time than the naive kernel. Due to this reason, SVE instruction based microkernel gives a noticeable performance enhancement of 14-16 Gflops than the naive microkernel.

**Conclusion:** Using SVE instructions (i.e. SIMD instructions) in microkernel gave a **performance enhancement of about 14 to 16 Gflops** for square matrix multiplication of different matrix sizes.

### Optimization 2: Choosing a micorkernel size


Commit Message	Commit Link
Added 4x4 microkernel	<a href="https://github.com/cse260-fa22/pa1-hgondali-shanant/h/commit/811619117ae7f8e82d2380070d7ecd5e5607e624">https://github.com/cse260-fa22/pa1-hgondali-shanant/h/commit/811619117ae7f8e82d2380070d7ecd5e5607e624</a>
Added 10x4 8x4 6x4 2x4 micorkernel	<a href="https://github.com/cse260-fa22/pa1-hgondali-shanant/h/commit/c63204394f742363e30f657ea33b02ffcc51a223">https://github.com/cse260-fa22/pa1-hgondali-shanant/h/commit/c63204394f742363e30f657ea33b02ffcc51a223</a>
Added 12x4 microkernel	<a href="https://github.com/cse260-fa22/pa1-hgondali-shanant/h/commit/c5c3ee8be237f1ef9df9a8d6d919713b1b600d7f">https://github.com/cse260-fa22/pa1-hgondali-shanant/h/commit/c5c3ee8be237f1ef9df9a8d6d919713b1b600d7f</a>
Added 14x4 and 16x4 micorkernel	<a href="https://github.com/cse260-fa22/pa1-hgondali-shanant/h/commit/da439cbfdaf76066826f8ab1dd5ae67645fad251">https://github.com/cse260-fa22/pa1-hgondali-shanant/h/commit/da439cbfdaf76066826f8ab1dd5ae67645fad251</a>
Added dumps for 12x4 14x4 16x4 kernel	<a href="https://github.com/cse260-fa22/pa1-hgondali-shanant/h/commit/d9227aa2e7046212ba7e485a83fedd856139c5cb">https://github.com/cse260-fa22/pa1-hgondali-shanant/h/commit/d9227aa2e7046212ba7e485a83fedd856139c5cb</a>

## Methodology and Reasoning:

Our microkernel is responsible for calculating multiplication results for  $M_r \times N_r$  block of matrix C. Throughout our experiment, we have kept  $N_r$  value as constant (i.e. equal to 4). This is with the insight that the vector length of Neoverse V1 is 256 bits [2]. Thus, by setting  $N_r$  value to 4, the four 64-bit double precision numbers will utilize vector width to its full capacity.

We could have not tried  $N_r$  values as 8, 12, 16, etc., but we reasoned out that there is more merit in using the available SVE registers for getting more depth along  $M_r$  dimension. This is because elements of  $M_r \times K_c$  block are going to be accessed more frequently and thus will be in L1 cache while the elements of  $K_c \times N_r$  subpanel are going to be accessed relatively less frequently and will be in the L2 cache.

Hence, in this experiment we implemented different microkernels by varying  $M_r$  size and keeping  $N_r$  as constant.



## Results Analysis:

For  $M_r$  value 2 to 14, we see an **increase** in performance for matrix multiplication for square matrices of size 2048. This is because register access costs least latency when compared to other forms of memory. Thus, as we increase  $M_r$  value, relatively more operands are read from registers which results in lesser time to fetch data. This is the reason for improvement in performance.

As we increase  $M_r$  value from 14 to 16, we see a **decrease** in performance. This is because SVE instruction set only has 32 scalable vector registers. When  $M_r$  value is set to 16, we run out of registers and thus need to access L1 cache to fetch some operand values. Since access to L1 cache causes more latency than access to registers, the performance for microkernel with  $M_r$  value as 16 is lesser than the microkernel with  $M_r$  value as 14.

We verified our reasoning by checking the assembly code for microkernels [1] with  $M_r$  value 12, 14, and 16. We clearly see that the microkernel with  $M_r$  value 16 has used up all 32 registers and **needs to use register z0.d twice (line 55 and 72 in 16x4.txt at [1]) for loading values**. Thus, in each iteration there will be one L1 cache access to get the required operand.


**Conclusion:** With all the results and analysis shared above, we chose **14 x 4 ( $M_r = 14$ ,  $N_r = 4$ ) as our micro-kernel size**

## Optimization 3: Choosing $M_c$ and $N_c$ values that are multiples of $M_r$ and $N_r$ respectively

Commit Message	Commit Link
----------------	-------------

Results for Mc value 70 and 64 for 14x4 microkernel	<a href="https://github.com/cse260-fa22/pa1-hgondali-shanant/commit/815be8656ee8b2ed61fe8464e444d35d1dd47048">https://github.com/cse260-fa22/pa1-hgondali-shanant/commit/815be8656ee8b2ed61fe8464e444d35d1dd47048</a>
Results for Nc value 64 and 70 for 14x4 microkernel	<a href="https://github.com/cse260-fa22/pa1-hgondali-shanant/commit/c2753b6059b82adab916e53afa9f2ffb117815b1">https://github.com/cse260-fa22/pa1-hgondali-shanant/commit/c2753b6059b82adab916e53afa9f2ffb117815b1</a>

**Methodology and Reasoning:** In our experiments, we came across this observation that we must choose Mc and Nc that are multiples of our Mr and Nr respectively (i.e. 14 and 4 respectively). This is because if we choose values that are not a multiple of Mr and Nr then each Mr x Kc and Kc x Nr subpanel will have padding and we will be performing unrequired computations in each subpanel. By choosing Mc and Nc that are multiples of our Mr and Nr respectively, we ensure that padding only needs to be done in subpanels that fall along the bottommost rows and rightmost columns of the matrix A and B.



**Results Analysis:** We notice a performance **difference of around 0.1 to 2 Gflops** between experiments where Mc and Nc values are multiples of Mr and Nr respectively and the case where they are not a multiple of Mr and Nr respectively. Due to space constraints, only plot for Mc value variation is shown.

**Conclusion:** The chosen value of **Mc and Nc must be a multiple of Mr and Nr** respectively.

#### Optimization 4: Searching for best possible (fine-tuned) values of Mc, Nc, and Kc



Commit Message	Commit Link
Results while varying Kc	<a href="https://github.com/cse260-fa22/pa1-hgondali-shanant/commit/3c0da5dfebb711f1256505fd1659db5abad2543b">https://github.com/cse260-fa22/pa1-hgondali-shanant/commit/3c0da5dfebb711f1256505fd1659db5abad2543b</a>
Results while varying Mc	<a href="https://github.com/cse260-fa22/pa1-hgondali-shanant/commit/d135ea15672e7636647d8686d5b3b655f48cf3d2">https://github.com/cse260-fa22/pa1-hgondali-shanant/commit/d135ea15672e7636647d8686d5b3b655f48cf3d2</a>
Results while varying Nc	<a href="https://github.com/cse260-fa22/pa1-hgondali-shanant/commit/8b0dce5f2f54343f52829c80a3494b5d9d0f040d">https://github.com/cse260-fa22/pa1-hgondali-shanant/commit/8b0dce5f2f54343f52829c80a3494b5d9d0f040d</a>


Perf counters while varying Mc, Nc, Kc

<https://github.com/cse260-fa22/pa1-hgondali-shanant/commit/ee08f528485990591a6be11f0d0112a3a7c9361a>

**Methodology and Reasoning:** Given the paucity of time, we figured out the best set of Kc, Mc, and Nc values for a 14 x 4 (Mr x Nr) size microkernel in the following manner:

1. Kept Mc = 70 (multiple of 14, i.e. Mr value), Nc = 64 (default value), and varied Kc value from 64 to 1024 in increments of 64. Found that Kc value of 448 gives us good performance results and hence used it further in experiments to find the best Mc and Nc value.
2. Kept Kc = 448, Nc = 64 (default value), and varied Mc value from 112 to 1232 in increments of 112 (multiple of 14, i.e. Mr value). Found that Mc value of 784 gives us good performance results and hence used it further in experiments to find the best Nc value.
3. Kept Mc = 784, Kc = 448, and varied Nc value from 64 to 704 in increments of 64. Found that Nc value of 128 gives us good performance results.





### Results Analysis:

**Varying Kc value:** The hardware brings entire cachelines worth of data to L1 cache when any one of its data element is accessed. Moreover, hardware and software prefetchers also bring adjacent cache lines into cache to take advantage of spatial locality.

Hence, if we do not have a larger Kc value then we would not be capitalizing upon the data (in the same cacheline and one in the prefetched cacheline) that is brought into cache. On the other hand, by having larger Kc value, our performance will improve because more number of data accesses done in the  $M_r \times K_c$  subpanel will be satisfied in the L1 cache itself. Our measurement results from the Linux Perf Tool [4] validate our reasoning by revealing that the **L1d cache miss rate decreases by 1.65%** as we move from Kc value 64 to 448.

Secondly, product of 448 (Kc), 14 ( $M_r$ ), 8 (double-precision size) roughly equals to 50 KB which can be accommodate entirely in the 64 KB L1-d cache.

Kc Value	L1-d Cache Miss Rate
64	3.65
448	2.00

**Varying Mc value:** As we increase the value of  $M_c$ , more number of  $M_r \times K_c$  panels could fit into the  $M_c \times K_c$  panel. As per the access pattern,  $M_c \times K_c$  panel resides in L3 cache and by increasing  $M_c$  size, more number of  $M_r \times K_c$  subpanel could be accommodated in the panel. Hence, when these subpanels are packed together, we can again take advantage of the cacheline worth of data that is brought in altogether while performing operations. The cost of bring  $M_c \times K_c$  panel into L3 cache is amortized over multiple  $M_r \times K_c$  subpanels that settle the reads for operations at the L1 cache itself. This leads to an increase in performance. Our measurement results from the Linux Perf Tool [4] validate our reasoning by revealing that the **L1d cache miss rate decreases by 0.56%** as we move from  $M_c$  value 112 to 784.

Mc Value	L1-d Cache Miss Rate
112	1.59
784	1.03

**Varying Nc value:** Increasing Nc value did not give us any significant improvement in performance. We suspect this is because increasing Nc value increases the number of Kc x Nr subpanels that could fit into Kc x Nc panel. Since Kc x Nr subpanel is accessed less frequently and is always pushed down to L2 cache as per the access pattern, there is no significant performance improvement by varying Nc.

**Conclusion:** Mc = 784, Kc = 448, Nc = 128 values gives us best possible performance enhancement. Finally, with all the four optimizations shared above, we are able to reach an average performance of ~21 Gflops.

#### Optimization 5: Using alternative SVE instructions

Commit Message	Commit Link
Using svmla_f64_x instruction in microkernel	<a href="https://github.com/cse260-fa22/pa1-hgondali-shanant/commit/7177bea95efef2725d94eca6f88df954bc7d67f9">https://github.com/cse260-fa22/pa1-hgondali-shanant/commit/7177bea95efef2725d94eca6f88df954bc7d67f9</a>
Using svmla_f64_z instruction in microkernel	<a href="https://github.com/cse260-fa22/pa1-hgondali-shanant/commit/7177bea95efef2725d94eca6f88df954bc7d67f9">https://github.com/cse260-fa22/pa1-hgondali-shanant/commit/7177bea95efef2725d94eca6f88df954bc7d67f9</a>

Instead of svmla\_f64\_m, we tried using svmla\_f64\_x and svmla\_f64\_z instructions to perform multiply, load, and add operations. This did not give any performance enhancement. Rather, implementation using svmla\_f64\_z instruction **decreased the performance by around 1-2 Gflops**. This is because svmla\_f64\_z instruction cost more CPU cycles than the other two instructions.

#### Optimization 6: Rearranging instructions within micro-kernel

Commit Message	Commit Link
Access B followed by A	<a href="https://github.com/cse260-fa22/pa1-hgondali-shanant/commit/db3e1de2bd7bc47448e614998e019427d26695a0">https://github.com/cse260-fa22/pa1-hgondali-shanant/commit/db3e1de2bd7bc47448e614998e019427d26695a0</a>

Inside the microkernel, we tried loading all values from matrix A first followed by B, and in other approach tried loading value of matrix B first followed by loading all values of matrix A. This did not give any performance enhancement because objdump revealed that compiler generated same assembly code for both the approaches.

#### Optimization 7: Using addition in place of multiplication to compute load address

Commit Message	Commit Link
Use add in place of multiply to compute load address	<a href="https://github.com/cse260-fa22/pa1-hgondali-shanant/commit/9a0a82ff927371f334a35c9e32c1c88ab727c774">https://github.com/cse260-fa22/pa1-hgondali-shanant/commit/9a0a82ff927371f334a35c9e32c1c88ab727c774</a>

Instead of svld1\_f64(npred, C+(ldc\*3)), we tried using svld1\_f64(npred, C+(ldc+ldc+ldc)) for all such load instructions. This did not give any performance enhancement because objdump revealed that compiler generated same assembly code for both the approaches.

#### Q2c Explanation of irregularities seen in graph shown in Q1b

Our implementation's performance increases steeply as we move from matrix size 32 to 512. This is because smaller the matrix size (lesser computations), the lesser we can take advantage of cacheline worth of data (64B) that has been brought into cache when any one of the data stored in cacheline is read. As a result, we are not amortizing the cost of data packing (bring data to caches) over enough computations and hence performance is lower.



Our implementation performs noticeably better than naive approach over all matrix sizes due to the data packing and SIMD-based vectorization approach that amortizes the cost of data access over multiple computations.

### **Q2d Analysis of Cache Behavior and Difference in Implementation from BLISLab Tutorial**

**Note:** Graphs, Performance Counter Readings, and Calculations referenced in this subsection are available in “Optimization 4” subsection in Q2b

We observe that increasing the value of  $K_c$  helps us take advantage of the cacheline worth of data that is brought into L1 cache for  $M_r \times K_c$  subpanel. As  $K_c$  value increases, we are able to amortize the cost of bringing a cacheline worth of data over more number of computations and hence we see an increase in performance. The performance counters showed clear decrease in L1 cache miss rate with increasing  $K_c$  because more cachelines (including prefetched ones) are used in larger size  $M_r \times K_c$  subpanel. The calculations and graph for average performance while varying  $K_c$  clearly shows that when  $K_c$  is value is more than 448 we are almost at the full capacity of L1 cache (64KB) and hence increasing  $K_c$  beyond that value leads to a decrease in performance.

Increasing  $M_c$  value helps us accommodate more  $M_r \times K_c$  subpanels in a single  $M_c \times K_c$  panel. Hence, cost of bringing  $M_c \times K_c$  panel into L3 cache is amortized over multiple  $M_r \times K_c$  subpanels that handle the reads for operations at the L1 cache itself. Performance counter readings clearly show a decrease in L1-d cache miss rate when  $M_c$  value is increased.

Increasing  $N_c$  value did not give us any significant improvement in performance because  $K_c \times N_r$  subpanel is accessed less frequently and is always pushed down to L2 cache as per the access pattern. Thus, there is no significant performance improvement by varying  $N_c$  and thereby fitting in more  $K_c \times N_r$  subpanels in  $K_c \times N_c$  panel.

Our implementation is based on row-major storage order while BLISLab tutorial is based on column-major storage order. For a column major order based implementation, its better to put  $K_c \times N_c$  subpanels of B into L3 cache (contiguous addresses within columns) rather than  $M_c \times K_c$  subpanel of A (non-contiguous addresses within rows). Similarly, for a row-major order based implementation, its better to put  $M_c \times K_c$  subpanel of A (contiguous addresses within rows) rather than  $K_c \times N_c$  subpanels of B into L3 cache (non-contiguous addresses within columns). Based on this, to take advantage of storage order during packing and further computation, arrangement of loops around micokernel differs in our and BLISLab implementation.

### **Q2e Future Work**

1. If we had more time, then we would like to implement the butterfly multiplication method described in BLISLab tutorial and see how much performance enhancement that can bring.
2. Our performance results are around 0.3 to 0.9 Gflops less than that of OpenBLAS implementation for different matrix sizes. We would like to do a more fine-grained parameter tuning to see if we can achieve better performance.

### **Q3 References**

[1] Assembly Code for 12x4, 14x4, and 16x4 microkernels.

<https://github.com/cse260-fa22/pa1-hgondali-shananth/tree/main/dumps>

[2] Neoverse V1 Specification

[https://en.wikichip.org/wiki/arm\\_holdings/microarchitectures/neoverse\\_v1](https://en.wikichip.org/wiki/arm_holdings/microarchitectures/neoverse_v1)

[3] SVE Architecture Fundamentals

<https://developer.arm.com/documentation/102476/0100/SVE-architecture-fundamentals>

[4] Linux Perf Tool

[https://perf.wiki.kernel.org/index.php/Main\\_Page](https://perf.wiki.kernel.org/index.php/Main_Page)

[5] Intel Intrinsics Guide

<https://www.intel.com/content/www/us/en/docs/intrinsics-guide/index.html>

#### **Q4 Extra Credit Work - Matrix Multiplication using AVX2 intrinsics**

Commit Message	Commit Link
Compared our implementation with OpenBLAS and Naive	<a href="https://github.com/cse260-fa22/pa1-hgondali-shanant/commit/0bc4029788597fb265be560df30714f2bbcf4f33">https://github.com/cse260-fa22/pa1-hgondali-shanant/commit/0bc4029788597fb265be560df30714f2bbcf4f33</a>

We implemented the matrix multiplication using AVX2 intrinsics on a x86 Ubuntu machine on Amazon EC2, and made use of the t2.micro instance type. The Makefile was modified to replace the SVE and ARM related flags with x86 and AVX2 flags. We modified the microkernel to make use of AVX2 intrinsics instead of SVE.

In order to optimize the use of registers, we use a 7x4 microkernel that makes use of exactly 16 registers. We experimented with microkernels of size 2x4, 4x4 and 8x4. There was a gradual performance improvement as our Mr value increased from 2 to 7 followed by a sudden drop when Mr is 8. This is because we have completely exhausted the registers when Mr is 7 which then causes data to be picked up from the L1 cache when Mr increases to 8.

We also experimented with the values of Kc, Mc and Nc and found that Kc = 224, Mc = 700, Nc = 120 to give us the best result. This can be attributed to cache hierarchy - the sizes of L3, L2 and L1 caches.

Unlike SVE, AVX2 does not make use of predicates and so our initial implementation worked only when the matrix size is a power of 2. To overcome this, we made use of a mask that changes according to the size of the leading dimension of the second matrix (Matrix B) that is input to the microkernel. Subsequently, a different form of the AVX2 intrinsic was chosen to perform masked loads and stores. The table below shows a comparison of the performances of the naive, OpenBLAS and our implementation of matrix multiplication. We achieve an average performance of **22.2 Gflops** for matrix sizes greater than 512 - which is around **65%** of the OpenBLAS implementation.

Matrix Size	Naive	OpenBLAS	Peak GF achieved in our Implementation
513	1.3525	31.805	22.01
1023	1.04145	34.1	22.135
1024	0.12865	34.75	22.65
1025	0.9514	33.645	22.075
2047	NA	35.17	22.51
2048	NA	35.125	21.78